

# SuayLang Research Plan (v0.2)

\*\*Build:\*\* `make research-pdf` regenerates [docs/RESEARCH\_PLAN.pdf](RESEARCH\_PLAN.pdf).

## Problem Statement

Many small languages ship a reference interpreter and a faster backend (VM/bytecode), but the two often diverge in subtle ways (semantics, diagnostics, determinism). SuayLang's goal is to make control flow explicit \*and\* make backend equivalence and diagnostic stability measurable and reproducible.

## Research Questions

- 1) Can an expression-oriented language with explicit control-flow operators (`dispatch`, `cycle`) be specified and tested with a small, committee-reviewable contract?
- 2) Can we provide strong, reproducible evidence that two implementations (interpreter vs. bytecode VM) are observationally equivalent over a stated v1 scope?

## Hypothesis (falsifiable)

Within the v1 scope, the interpreter and the VM are observationally equivalent under a fixed observation policy (termination class, normalized stdout, returned value when comparable, and error kind+span).

## Method (how we will test it)

- bullet Semantic contract + golden diagnostics\*\*: a small v1 contract document plus a suite of valid/invalid programs with golden snapshots for error kind/span/message shape.
- bullet Differential testing (interpreter vs. VM)\*\*: deterministic generator by seed, multi-seed runs, size buckets (S/M/L), timeouts.
- bullet Comparator + normalization\*\*: normalize paths and whitespace; compare outputs exactly after normalization; compare errors by kind + span (and stable formatting).
- bullet Minimization\*\*: when a divergence is found, shrink the program and store a minimized regression case with metadata.
- bullet Coverage reporting\*\*: record feature coverage (AST nodes and/or opcode families) to show exploration breadth.
- bullet Benchmarks\*\*: measure parse/compile/interp/VM times with warmups and repeats; report median and p90 with raw samples.

## Metrics & Success Criteria

All metrics are produced by repository commands and saved under `results/`.

\*\*Equivalence (primary)\*\*

bullet CI mode: seeds 0 .. 9, N=500 / seed  $\Rightarrow$  \*\*divergences = 0\*\*.

bullet Full mode: seeds 0 .. 99, N=2000 / seed  $\Rightarrow$  \*\*divergences = 0\*\*.

\*\*Diagnostics stability (primary)\*\*

bullet For the invalid-program contract corpus: \*\*100% match\*\* on error kind + span, and stable message prefix/shape.

\*\*Coverage (supporting)\*\*

bullet Coverage report includes counts by feature class (dispatch/cycle/functions/collections/errors). Success criterion: no major feature class is zero-covered in CI mode.

\*\*Benchmarks (supporting; not a correctness proof)\*\*

bullet Report median and p90 for parse, compile-to-bytecode, interpreter runtime, and VM runtime with  $\geq 20$  repeats + warmup. Success: the benchmark runner emits raw samples and environment metadata.

## Experimental Protocol (repeatability)

bullet Runs are deterministic by \*\*seed\*\* and generator configuration.

bullet Profiles:

bullet<sub>CI</sub>: fast gate (seeds 0 .. 9, smaller N).

bullet<sub>Full</sub>: long local run (seeds 0 .. 99, larger N).

bullet Every run captures:

bullet Python version, OS, CPU info (best-effort), and git commit hash.

bullet Per-seed breakdown and size-bucket breakdown.

bullet Artifacts:

bullet<sub>results/diff\_report.json + results/diff\_report.md</sub>

bullet<sub>results/coverage.json + results/coverage.md</sub>

bullet<sub>results/bench\_raw.json + results/benchmarks.md</sub>

## Threats to Validity

bullet<sub>Generator bias</sub>: random generation may over/under-sample important features.

bullet<sub>Observation policy limitations</sub>: “equivalence” is with respect to the chosen observable outcomes.

bullet<sub>Timeouts</sub>: timeouts may hide non-termination differences.

bullet<sub>Host effects</sub>: Python runtime and OS scheduling noise affects benchmark timing.

bullet<sub>Scope gaps</sub>: behavior outside the declared v1 subset is intentionally unspecified.

## Expected Results + What Would Falsify the Hypothesis

\*\*Expected:\*\* no divergences within v1 scope under CI and full differential testing profiles; stable diagnostics for the contract corpus.

\*\*Falsifiers:\*\*

bullet Any reproducible divergence in stdout/value/error kind+span for the same program input between interpreter and VM.

bullet Any change that breaks golden diagnostic snapshots for the v1 contract invalid-program corpus (without an explicit, documented version bump).

## Timeline (milestones)

- 1) Finalize v1 contract and golden diagnostics corpus.
- 2) Scale differential testing to multi-seed CI + full profiles; enable minimization.
- 3) Add feature coverage reporting (AST/opcode) and publish reports to `results/`.
- 4) Add benchmark suite + runner with raw samples and noise controls.
- 5) Tooling polish: one-command install + `suay` CLI closed loop.

## References

- bullet Plotkin (1981), \*A Structural Approach to Operational Semantics\*.
- bullet Kahn (1987), \*Natural Semantics\*.
- bullet McKeeman (1998), \*Differential Testing for Software\*.
- bullet Claessen & Hughes (2000), \*QuickCheck\*.
- bullet Wang et al. (2011), \*Finding and Understanding Bugs in C Compilers\*.
- bullet De et al. (2014), \*Compiler Validation via Equivalence Modulo Inputs\*.
- bullet Maranget (2008), \*Compiling Pattern Matching to Good Decision Trees\*.